

Upload Kaggle dataset in google colab:

```
MUSIC.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

Code + Text

[1] pip install -g kaggle

(module) colab
▶ from google.colab import files
  files.upload()

Choose Files kaggle.json
• kaggle.json(application/json) - 65 bytes, last modified: 5/4/2022 - 100% done
  Saving kaggle.json to kaggle (2).json
  {'kaggle.json': b'{"username": "shanjai34", "key": "2f1c00526781fb25f9297e6361576333"}'}
```

```
[3] !mkdir ~/.kaggle
    !cp kaggle.json ~/.kaggle/

mkdir: cannot create directory '/root/.kaggle': File exists

[4] !chmod 600 ~/.kaggle/kaggle.json

[5] !kaggle datasets download andradaolteanu/gtzan-dataset-music-genre-classification
gtzan-dataset-music-genre-classification.zip: Skipping, found more recently modified local copy (use --force to force download)
```

Export data:

```
export data

[6] from zipfile import ZipFile
    file_name = "gtzan-dataset-music-genre-classification.zip"

    with ZipFile(file_name, "r") as zip:
        zip.extractall()
        print("Done")

Done
```

Importing libraries:

## ▼ import libraries

```
✓ 7 import numpy as np
4s import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
import sklearn.metrics as skm
import sklearn.model_selection as skms
import sklearn.preprocessing as skp
import random
import librosa, IPython
import librosa.display as lplt
seed = 12
np.random.seed(seed)
```

```
df = pd.read_csv('Data/features_3_sec.csv')
df.head()
```

	filename	length	chroma_stft_mean	chroma_stft_var	rms_mean	rms_var	spectral_centroid_mean	spectral_centroid_var	spectral_bandwidth_mean	spectral_bandwidth_var	...	m
0	blues.00000.0.wav	66149	0.335406	0.091048	0.130405	0.003521	1773.065032	167541.630869	1972.744388	117335.771563	...	
1	blues.00000.1.wav	66149	0.343065	0.086147	0.112699	0.001450	1816.693777	90525.690866	2010.051501	65671.875673	...	
2	blues.00000.2.wav	66149	0.346815	0.092243	0.132003	0.004620	1788.539719	111407.437613	2084.565132	75124.921716	...	
3	blues.00000.3.wav	66149	0.363639	0.086856	0.132565	0.002448	1655.289045	111952.284517	1960.039988	82913.639269	...	
4	blues.00000.4.wav	66149	0.335579	0.088129	0.143289	0.001701	1630.656199	79667.267654	1948.503884	60204.020268	...	

5 rows × 60 columns

✓  
0s



```
print("Dataset has",df.shape)
print("Count of Positive and Negative samples")
df.label.value_counts().reset_index()
```



Dataset has (9990, 60)  
Count of Positive and Negative samples

	index	label
0	blues	1000
1	jazz	1000
2	metal	1000
3	pop	1000
4	reggae	1000
5	disco	999
6	classical	998
7	hiphop	998
8	rock	998
9	country	997



✓  
0s


```
[10] audio_fp = 'Data/genres_original/blues/blues.00000.wav'
      audio_data, sr = librosa.load(audio_fp)
      audio_data, _ = librosa.effects.trim(audio_data)
```

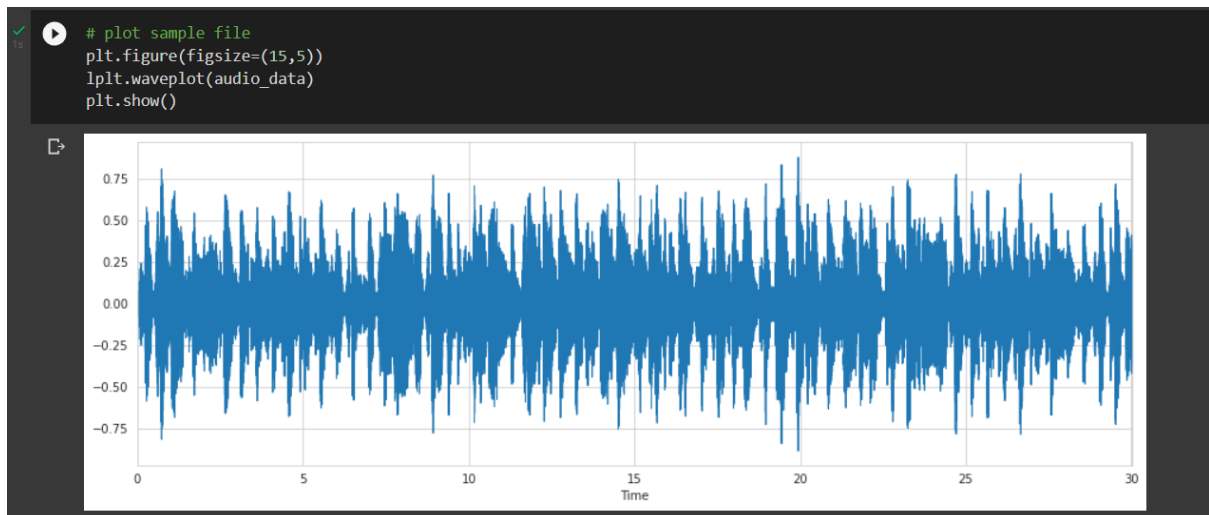
✓  
4s

```
[11] # play sample file
      IPython.display.Audio(audio_data, rate=sr)
```



0:00 / 0:30





Preprocessing:

## PCA

```
data = df.iloc[0:, 1:]
y = data['label']
X = data.loc[:, data.columns != 'label']

# normalize
cols = X.columns
min_max_scaler = sklearn.preprocessing.MinMaxScaler()
np_scaled = min_max_scaler.fit_transform(X)
X = pd.DataFrame(np_scaled, columns = cols)


# Top 2 pca components
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
principalComponents = pca.fit_transform(X)
principalDf = pd.DataFrame(data = principalComponents, columns = ['pc1', 'pc2'])

# concatenate with target label
finalDf = pd.concat([principalDf, y], axis = 1)

plt.figure(figsize = (16, 9))
sns.scatterplot(x = "pc1", y = "pc2", data = finalDf, hue = "label", alpha = 0.7, s = 100);
```

```
plt.title('PCA on Genres', fontsize = 20)
plt.xticks(fontsize = 14)
plt.yticks(fontsize = 10);
plt.xlabel("Principal Component 1", fontsize = 15)
plt.ylabel("Principal Component 2", fontsize = 15)
plt.savefig("PCA_Scattert.png")
```



Feature Extraction:

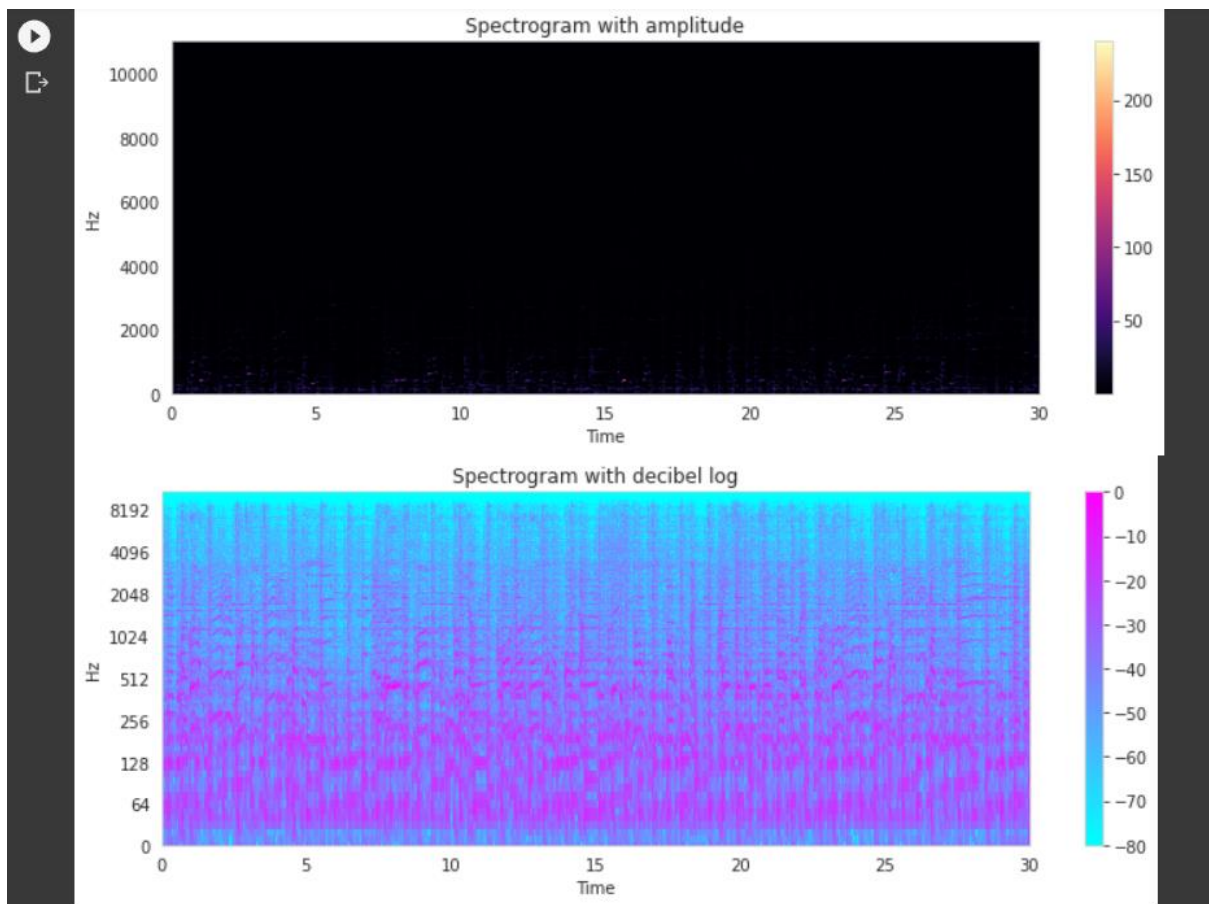
## FEATURE EXTRACTION

```
[14] # Default FFT window size
n_fft = 2048 # window size
hop_length = 512 # window hop length for STFT


stft = librosa.stft(audio_data, n_fft=n_fft, hop_length=hop_length)
stft_db = librosa.amplitude_to_db(stft, ref=np.max)

plt.figure(figsize=(12,4))
lplt.specshow(stft, sr=sr, x_axis='time', y_axis='hz')
plt.colorbar()
plt.title("Spectrogram with amplitude")
plt.show()

plt.figure(figsize=(12,4))
lplt.specshow(stft_db, sr=sr, x_axis='time', y_axis='log', cmap='cool')
plt.colorbar()
plt.title("Spectrogram with decibel log")
plt.show()
```



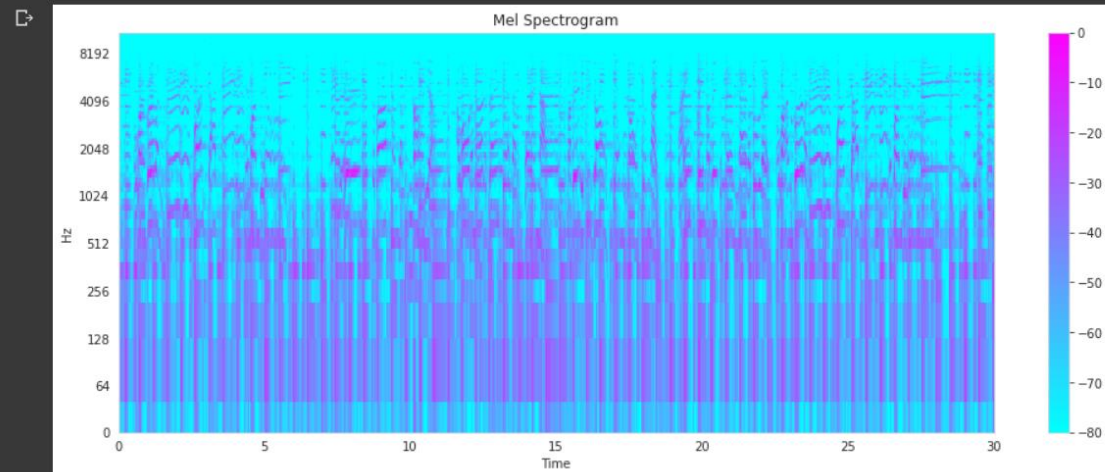
```
# plot zoomed audio wave
start = 1000
end = 1200
plt.figure(figsize=(16,4))
plt.plot(audio_data[start:end])
plt.show()
```



```

mel_spec = librosa.feature.melspectrogram(audio_data, sr=sr)
mel_spec_db = librosa.amplitude_to_db(mel_spec, ref=np.max)
plt.figure(figsize=(16,6))
lplt.specshow(mel_spec_db, sr=sr, hop_length=hop_length, x_axis='time', y_axis='log', cmap='cool')
plt.colorbar()
plt.title("Mel Spectrogram")
plt.show()


```



```

chroma = librosa.feature.chroma_stft(audio_data, sr=sr)
plt.figure(figsize=(16,6))
lplt.specshow(chroma, sr=sr, x_axis='time', y_axis='chroma', cmap='coolwarm')
plt.colorbar()
plt.title("Chroma Features")
plt.show()

```




```

x = df[["label", "tempo"]]

fig, ax = plt.subplots(figsize=(16, 8));
sns.boxplot(x = "label", y = "tempo", data = x, palette = 'husl');

plt.title('BPM Boxplot for Genres', fontsize = 20)
plt.xticks(fontsize = 14)
plt.yticks(fontsize = 10);
plt.xlabel("Genre", fontsize = 15)
plt.ylabel("BPM", fontsize = 15)
plt.savefig("BPM_Boxplot.png")

```



## ▼ Missing value treatment

```

[20] # find all columns with any NA values
print("Columns with NA values are", list(df.columns[df.isnull().any()]))

Columns with NA values are []

```

## ▼ Encode genre label

```

[21] # map labels to index
label_index = dict()
index_label = dict()
for i, x in enumerate(df.label.unique()):
    label_index[x] = i
    index_label[i] = x
print(label_index)
print(index_label)

{'blues': 0, 'classical': 1, 'country': 2, 'disco': 3, 'hiphop': 4, 'jazz': 5, 'metal': 6, 'pop': 7, 'reggae': 8, 'rock': 9}
{0: 'blues', 1: 'classical', 2: 'country', 3: 'disco', 4: 'hiphop', 5: 'jazz', 6: 'metal', 7: 'pop', 8: 'reggae', 9: 'rock'}

```



## Train and Test split:

```
[22] # update labels in df to index
df.label = [label_index[l] for l in df.label]
```

Train&Test

```
[23] # shuffle samples
df_shuffle = df.sample(frac=1, random_state=seed).reset_index(drop=True)
```

```
[24] # remove irrelevant columns
df_shuffle.drop(['filename', 'length'], axis=1, inplace=True)
df_y = df_shuffle.pop('label')
df_X = df_shuffle

# split into train dev and test
X_train, df_test_valid_X, y_train, df_test_valid_y = skms.train_test_split(df_X, df_y, train_size=0.7, random_state=seed, stratify=df_y)
X_dev, X_test, y_dev, y_test = skms.train_test_split(df_test_valid_X, df_test_valid_y, train_size=0.66, random_state=seed, stratify=df_test_valid_y)
```

```
print(f"Train set has {X_train.shape[0]} records out of {len(df_shuffle)} which is {round(X_train.shape[0]/len(df_shuffle)*100)}%")
print(f"Dev set has {X_dev.shape[0]} records out of {len(df_shuffle)} which is {round(X_dev.shape[0]/len(df_shuffle)*100)}%")
print(f"Test set has {X_test.shape[0]} records out of {len(df_shuffle)} which is {round(X_test.shape[0]/len(df_shuffle)*100)}%")
```

Train set has 6993 records out of 9990 which is 70%  
Dev set has 1978 records out of 9990 which is 20%  
Test set has 1019 records out of 9990 which is 10%

```
[27] scaler = skp.StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.columns)
X_dev = pd.DataFrame(scaler.transform(X_dev), columns=X_train.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns=X_train.columns)
```

```
[28] import tensorflow as tf
print("TF version:-", tf.__version__)
import keras as k
tf.random.set_seed(seed)
```

TF version:- 2.8.0

```
ACCURACY_THRESHOLD = 0.94

class myCallback(k.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        if(logs.get('val_accuracy') > ACCURACY_THRESHOLD):
            print("\n\nStopping training as we have reached %2.2f%% accuracy!" %(ACCURACY_THRESHOLD*100))
            self.model.stop_training = True

def trainModel(model, epochs, optimizer):
    batch_size = 128
    callback = myCallback()
    model.compile(optimizer=optimizer,
                  loss='sparse_categorical_crossentropy',
                  metrics='accuracy'
    )
    return model.fit(X_train, y_train, validation_data=(X_dev, y_dev), epochs=epochs,
                    batch_size=batch_size, callbacks=[callback])

def plotHistory(history):
    print("Max. Validation Accuracy", max(history.history["val_accuracy"]))
    pd.DataFrame(history.history).plot(figsize=(12,6))
    plt.show()
```

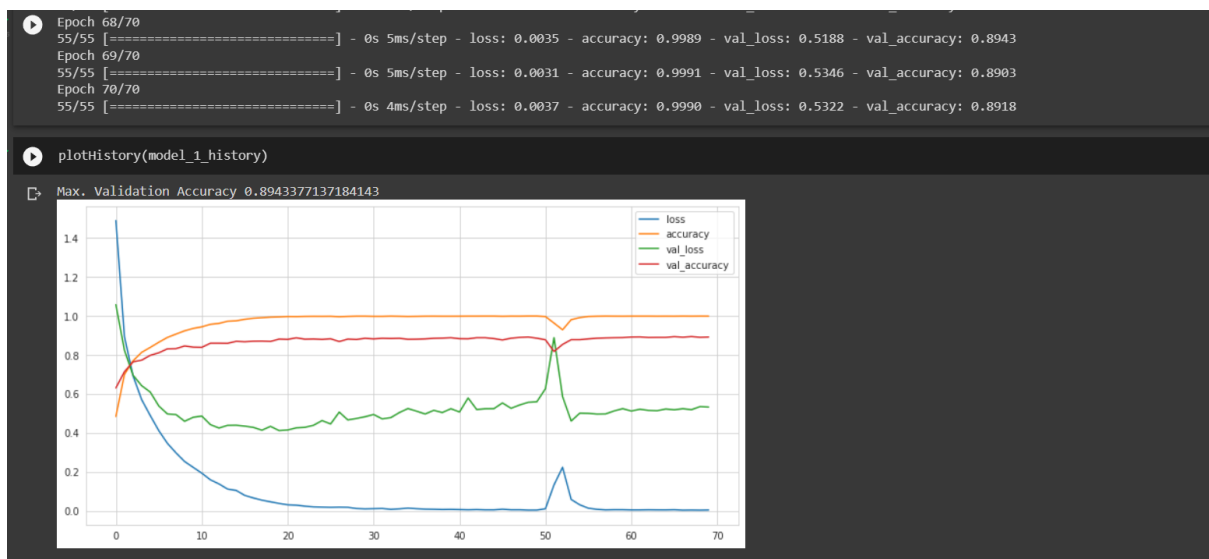
## Training:

```
18s model_1 = k.models.Sequential([
    k.layers.Dense(256, activation='relu', input_shape=(X_train.shape[1],)),
    k.layers.Dense(128, activation='relu'),
    k.layers.Dense(64, activation='relu'),
    k.layers.Dense(10, activation='softmax'),
])
print(model_1.summary())
model_1_history = trainModel(model=model_1, epochs=70, optimizer='adam')
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	14848
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 10)	650

=====  
Total params: 56,650  
Trainable params: 56,650  
Non-trainable params: 0





```
model_2 = k.models.Sequential([
    k.layers.Dense(512, activation='relu', input_shape=(X_train.shape[1],)),
    k.layers.Dropout(0.2),

    k.layers.Dense(256, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(128, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(64, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(10, activation='softmax'),
])
print(model_2.summary())
model_2_history = trainModel(model=model_2, epochs=100, optimizer='adam')
```



Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 512)	29696
dropout (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 256)	131328
dropout_1 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 10)	650


```

Epoch 97/100
55/55 [=====] - 1s 16ms/step - loss: 0.0356 - accuracy: 0.9888 - val_loss: 0.3379 - val_accuracy: 0.9257
Epoch 98/100
55/55 [=====] - 1s 13ms/step - loss: 0.0316 - accuracy: 0.9896 - val_loss: 0.4207 - val_accuracy: 0.9146
Epoch 99/100
55/55 [=====] - 1s 18ms/step - loss: 0.0393 - accuracy: 0.9873 - val_loss: 0.4001 - val_accuracy: 0.9146
Epoch 100/100
55/55 [=====] - 1s 16ms/step - loss: 0.0383 - accuracy: 0.9880 - val_loss: 0.3295 - val_accuracy: 0.9262

```

```
[33] plotHistory(model_2_history)
```

Max. Validation Accuracy 0.9312436580657959



```

model_3 = k.models.Sequential([
    k.layers.Dense(512, activation='relu', input_shape=(X_train.shape[1],)),
    k.layers.Dropout(0.2),

    k.layers.Dense(256, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(128, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(64, activation='relu'),
    k.layers.Dropout(0.2),

    k.layers.Dense(10, activation='softmax'),
])
print(model_3.summary())
model_3_history = trainModel(model=model_3, epochs=700, optimizer='sgd')


```

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 512)	29696
dropout_4 (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 256)	131328
dropout_5 (Dropout)	(None, 256)	0
dense_11 (Dense)	(None, 128)	32896
dropout_6 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 64)	8256
dropout_7 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 10)	650

```
55/55 [=====] - 1s 10ms/step - loss: 0.0421 - accuracy: 0.9866 - val_loss: 0.3340 - val_accuracy: 0.9151
Epoch 696/700
55/55 [=====] - 1s 9ms/step - loss: 0.0398 - accuracy: 0.9864 - val_loss: 0.3339 - val_accuracy: 0.9151
Epoch 697/700
55/55 [=====] - 1s 10ms/step - loss: 0.0396 - accuracy: 0.9890 - val_loss: 0.3327 - val_accuracy: 0.9125
Epoch 698/700
55/55 [=====] - 1s 10ms/step - loss: 0.0450 - accuracy: 0.9871 - val_loss: 0.3387 - val_accuracy: 0.9130
Epoch 699/700
55/55 [=====] - 1s 10ms/step - loss: 0.0436 - accuracy: 0.9868 - val_loss: 0.3358 - val_accuracy: 0.9151
Epoch 700/700
55/55 [=====] - 1s 9ms/step - loss: 0.0435 - accuracy: 0.9873 - val_loss: 0.3320 - val_accuracy: 0.9151
```

plotHistory(model\_3\_history)

Max. Validation Accuracy 0.9186046719551086



```
model_4 = k.models.Sequential([
    k.layers.Dense(1024, activation='relu', input_shape=(X_train.shape[1],)),
    k.layers.Dropout(0.3),

    k.layers.Dense(512, activation='relu'),
    k.layers.Dropout(0.3),

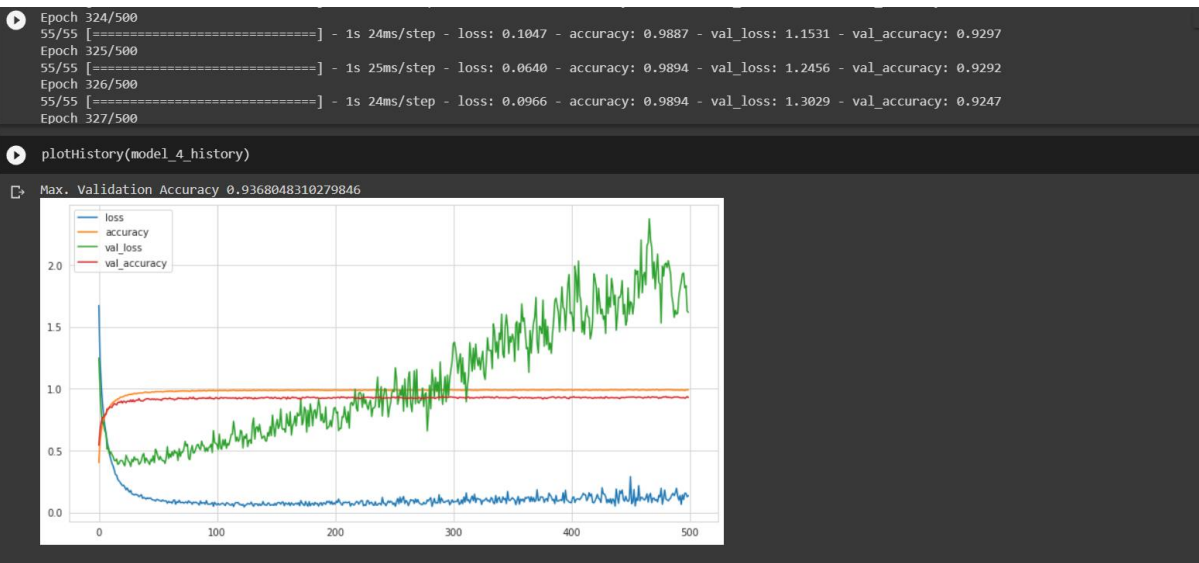
    k.layers.Dense(256, activation='relu'),
    k.layers.Dropout(0.3),

    k.layers.Dense(128, activation='relu'),
    k.layers.Dropout(0.3),

    k.layers.Dense(64, activation='relu'),
    k.layers.Dropout(0.3),

    k.layers.Dense(10, activation='softmax'),
])
print(model_4.summary())
model_4_history = trainModel(model=model_4, epochs=500, optimizer='rmsprop')
=====
```

dense_14 (Dense)	(None, 1024)	59392
dropout_8 (Dropout)	(None, 1024)	0
dense_15 (Dense)	(None, 512)	524800
dropout_9 (Dropout)	(None, 512)	0
dense_16 (Dense)	(None, 256)	131328
dropout_10 (Dropout)	(None, 256)	0
dense_17 (Dense)	(None, 128)	32896
dropout_11 (Dropout)	(None, 128)	0
dense_18 (Dense)	(None, 64)	8256
dropout_12 (Dropout)	(None, 64)	0
dense_19 (Dense)	(None, 10)	650



```
✓ [38] test_loss, test_acc = model_4.evaluate(X_test, y_test, batch_size=128)
28 print("The test Loss is :",test_loss)
    print("\nThe Best test Accuracy is :",test_acc*100)

8/8 [=====] - 0s 5ms/step - loss: 1.7462 - accuracy: 0.9293
The test Loss is : 1.7461915016174316

The Best test Accuracy is : 92.934250831604
```

Testing Prediction:

```
✓ ▶ dict2={0: 'blues', 1: 'classical', 2: 'country', 3: 'disco', 4: 'hiphop', 5: 'jazz', 6: 'metal', 7: 'pop', 8: 'reggae', 9: 'rock'}
    predicted_values=pd.DataFrame(y_test)
    predicted_values=predicted_values.replace({"label":dict2})
    predicted_values.reset_index(drop=True, inplace=True)
    predicted_values=list(predicted_values.label)
    i=0
    for item in predicted_values:
        print("The prediction of song no. {} in our dataset is {}".format(i,item))
        i=i+1

The prediction of song no. 957 in our dataset is rock
The prediction of song no. 958 in our dataset is metal
The prediction of song no. 959 in our dataset is classical
The prediction of song no. 960 in our dataset is blues
The prediction of song no. 961 in our dataset is jazz
The prediction of song no. 962 in our dataset is hiphop
The prediction of song no. 963 in our dataset is disco
The prediction of song no. 964 in our dataset is reggae
The prediction of song no. 965 in our dataset is blues
The prediction of song no. 966 in our dataset is rock
The prediction of song no. 967 in our dataset is pop
The prediction of song no. 968 in our dataset is blues
The prediction of song no. 969 in our dataset is rock
The prediction of song no. 970 in our dataset is disco
The prediction of song no. 971 in our dataset is pop
The prediction of song no. 972 in our dataset is classical
The prediction of song no. 973 in our dataset is blues
The prediction of song no. 974 in our dataset is rock
The prediction of song no. 975 in our dataset is reggae
The prediction of song no. 976 in our dataset is hiphop
```

